

Cross-Modal Audio-Visual Deepfake Detection for Biometric Security

Yash Khairnar
San Jose State University
San Jose, California, USA
yash.khairnar@sjsu.edu

Abstract

Deepfake synthesis has reached production quality sufficient to undermine face- and voice-based biometric authentication. Checkpoint 1 established a ResNet-50 visual baseline achieving 80.07% accuracy with FAR = 15.79% and EER = 18.95% on the DeepFake Detection Challenge (DFDC) dataset. Checkpoint 2 presents *AVConsist*, a cross-modal audio-visual deepfake detection framework that jointly encodes face frames via a Vision Transformer (ViT-Small/16) and voice via a lightweight CNN audio encoder, connected by a bidirectional cross-modal attention module that explicitly models audio-visual consistency. Trained and evaluated on the Localized Audio-Visual Deepfake (LAV-DF) dataset—a significantly harder benchmark featuring partial, content-driven manipulations across 136,304 clips—AVConsist achieves 88.93% accuracy, FAR = 9.20%, FRR = 12.93%, EER = 10.80%, and AUC = 0.9638 on a 10,000-clip stratified subset, outperforming the visual-only ResNet-50 baseline across all five biometric metrics. A data-scaling experiment across 300, 1,000, and 10,000 clips demonstrates a clear improvement trend (EER: 39.92% → 16.67% → 10.80%), confirming that transformer architectures benefit substantially from larger training corpora.

1 Refined Problem Statement and Motivation

1.1 Problem Definition

Biometric authentication systems that rely on face or voice verification are critically vulnerable to audio-visual deepfakes—attacks where an adversary simultaneously synthesizes photorealistic facial video and a cloned voice to impersonate a target individual. Formally, given a video clip $V = \{F_1, \dots, F_T, A\}$ of T face frames and audio track A , the goal is to learn a binary classifier $f_\theta : V \rightarrow \{0, 1\}$ where 0 denotes genuine and 1 denotes synthesized, generalizing to unseen manipulation methods and speaker identities.

1.2 Motivation

Existing detection methods operate on face *or* voice in isolation, remaining blind when only one modality is attacked. The few multi-modal approaches that exist are framed as media forensics problems and report generic metrics such as AUC and accuracy—none evaluate performance as a biometric security system using FAR, FRR, and EER at operationally meaningful thresholds. The ISO/IEC 19795-1 standard [8] provides this framework, which the deepfake detection community has largely not adopted. AVConsist explicitly bridges this gap.

2 System Architecture

2.1 Architecture Overview

Figure 1 presents the complete AVConsist pipeline comprising five functional blocks aligned with standard biometric system design.




Figure 1: AVConsist system architecture. Input video is split into face frames (ViT-Small branch, blue) and audio (CNN encoder branch, purple). Bidirectional cross-modal attention models audio-visual consistency before binary classification.

2.2 Component Description

(1) Input Interface: Raw .mp4 video clips from the LAV-DF dataset [11], comprising real clips and manipulated clips with partial audio and/or visual forgeries.

(2) **Preprocessing Pipeline:** Face frames are extracted at 16 evenly spaced intervals per clip and resized to 224×224 . Audio is extracted via `ffmpeg`, resampled to 16 kHz mono WAV, and converted to a 128-bin log-mel spectrogram over a 3-second window using `torchaudio`. ImageNet mean/std normalization is applied to face tensors.

(3) **Feature Extraction & Template Generation:** ViT-Small/16 produces face feature tensor $F_v \in \mathbb{R}^{B \times T \times 512}$, the CNN audio encoder produces voice feature vector $F_a \in \mathbb{R}^{B \times 512}$. Preprocessed embeddings are cached as .pt tensors for reuse.

(4) **Matcher/Classifier:** The cross-modal attention module—the architectural novelty—uses bidirectional MultiheadAttention (8 heads, dim 512) with LayerNorm and GELU feed-forward sublayers to model face-voice consistency. Fused features pass through $\text{Linear}(1024 \rightarrow 256) \rightarrow \text{ReLU} \rightarrow \text{Dropout}(0.3) \rightarrow \text{Linear}(256 \rightarrow 1) \rightarrow \text{Sigmoid}$ for binary classification.

(5) **Template Database Store:** Preprocessed .pt tensors are cached on Google Drive for reproducible evaluation and 1:1 identity comparison across sessions.

3 Advanced Methodology

3.1 Architecture Choices and Hyperparameters

ViT-Small/16 [13] is initialized from ImageNet-21k pretrained weights. Encoder weights are frozen for the first 5 warmup epochs to allow the cross-modal attention and classifier head to stabilize before full fine-tuning—a two-stage strategy that prevents catastrophic forgetting. The cross-modal attention module uses 8 attention heads with dimension 512, dropout rate 0.1, and GELU activations in the feed-forward sublayers. The ResNet-50 backbone used in our visual baseline follows He et al. [9].

3.2 Mathematical Formalism

Let $F_v \in \mathbb{R}^{B \times T \times d}$ be the face feature tensor and $F_a \in \mathbb{R}^{B \times d}$ the voice feature vector ($d = 512$). The cross-modal attention computes:

$$V_{\text{ctx}} = \text{LayerNorm}\left(F_v + \text{MHA}(F_v, \tilde{F}_a, \tilde{F}_a)\right) + \text{FFN}(\cdot) \quad (1)$$

$$A_{\text{ctx}} = \text{LayerNorm}\left(\tilde{F}_a + \text{MHA}(\tilde{F}_a, F_v, F_v)\right) + \text{FFN}(\cdot) \quad (2)$$

where $\tilde{F}_a = \text{expand}(F_a, T)$ broadcasts the voice vector over the temporal dimension. The fused representation is:

$$z = \left[\text{mean}_T(V_{\text{ctx}}); \text{mean}_T(A_{\text{ctx}}) \right] \in \mathbb{R}^{B \times 1024} \quad (3)$$

The loss function is weighted binary cross-entropy to handle class imbalance (our subset: 50% real, 50% fake, balanced by design):

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \left[w_+ y_i \log \sigma(\hat{z}_i) + (1 - y_i) \log(1 - \sigma(\hat{z}_i)) \right] \quad (4)$$

where $w_+ = n_{\text{real}}/n_{\text{fake}}$ is the positive class weight.

3.3 Training Strategy

- **Optimizer:** AdamW with weight decay 10^{-4} .

- **Learning rate:** 10^{-4} during warmup (epochs 1–5), decayed to 10^{-5} after encoder unfreeze (epochs 6–20) with cosine annealing.
- **Batch size:** 8 (constrained by T4 GPU VRAM with ViT active).
- **Gradient clipping:** Max norm = 1.0 to prevent exploding gradients.
- **Data augmentation:** Random horizontal flip and color jitter on face frames; SpecAugment time-masking and frequency-masking on log-mel spectrograms to prevent overfitting.
- **Dataset split:** Stratified 70/15/15 train/val/test.
- **Training data:** 10,000 clips (5,000 real + 5,000 fake) from LAV-DF [11].
- **Hardware:** NVIDIA T4 GPU (Google Colab).

4 Comprehensive Related Work

4.1 Unimodal Visual Detectors

FaceForensics++ (FF++) [2] established the dominant visual benchmark with 1,000 manipulated videos across four forgery methods, showing CNN-based detectors generalize poorly to unseen manipulations. Our ResNet-50 baseline (Checkpoint 1) follows this transfer-learning paradigm, achieving AUC = 0.8736 but FAR = 15.79%—confirming the inadequacy of visual-only detection for biometric security.

Li et al. [3] improved generalization via frequency-domain features, arguing that generative models leave spectral fingerprints invisible to spatial CNNs. We incorporate this insight in our audio branch: the CNN encoder operates on log-mel spectrograms that expose vocoder artifacts in the frequency domain. Unlike Li et al., our approach jointly models both visual and audio frequency signatures.

Liu et al. [1] conducted a systematic cross-dataset evaluation, finding models drop 20–30 percentage points on held-out generation methods. Our three-scale training experiment (300, 1,000, 10,000 clips) confirms this finding and shows that sufficient data volume can overcome distribution shift, with EER improving from 39.92% to 10.80% as training scale increases.

4.2 Audio-Visual Joint Detectors

AVoid-DF [4] pioneered joint face-audio modeling with late fusion, achieving state-of-the-art on FakeAVCeleb [10]. Our work differs by using cross-modal attention rather than late score fusion, enabling feature-level audio-visual consistency modeling rather than merely combining independent predictions.

AVFF [5] (CVPR 2024) introduced contrastive self-supervised pre-training on real audio-visual pairs, achieving AUC = 0.991 on FakeAVCeleb. AVConsist differs by not requiring a separate self-supervised pre-training stage, making it more practical for deployment with limited compute. Our AUC = 0.9638 on the harder LAV-DF benchmark is competitive given this architectural simplification.

ART-AVDF [6] (ACM MM 2024) adds adversarially robust training, showing that naive fusion is brittle against adaptive adversaries. While we do not yet implement adversarial training, our cross-modal attention architecture requires consistency across both streams simultaneously, providing inherent robustness to single-modality attacks.

AVMCD [7] uses ViViT + AST with a synchronization analysis module, reaching 96.1% accuracy on FakeAVCeleb. AVConsist uses ViT-Small (not ViViT) and adds explicit bidirectional cross-attention. The difference in benchmark datasets (FakeAVCeleb vs. LAV-DF) makes direct comparison difficult, but our AUC = 0.9638 on the harder LAV-DF benchmark is encouraging.

4.3 Biometric Framing Gap

A critical observation across all prior work is the absence of biometric-standard evaluation. None of [2]–[7] report FAR, FRR, or EER as primary metrics, and none analyze performance at operationally relevant FAR thresholds. The ISO/IEC 19795-1 standard [8] provides this framework. AVConsist is—to our knowledge—the first audio-visual deepfake detector to adopt ISO-aligned biometric evaluation as its primary reporting standard.

5 Preliminary Results and Comparison

5.1 Quantitative Results

Table 1 presents all biometric evaluation metrics. AVConsist outperforms the ResNet-50 baseline across all five metrics despite being evaluated on the significantly harder LAV-DF benchmark, which features partial content-driven manipulations rather than the whole-clip forgeries in DFDC.

Table 1: Baseline vs. AVConsist — biometric metrics. Datasets differ: DFDC (CP1) vs. LAV-DF (CP2). Red = security-critical FAR.

Metric	ResNet-50 (CP1)	AVConsist (CP2)
Accuracy	80.07%	88.93%
FAR	15.79%	9.20%
FRR	24.03%	12.93%
EER	18.95%	10.80%
AUC	0.8736	0.9638
TP / TN / FP / FN	117/128/24/37	653/681/69/97
Dataset	DFDC	LAV-DF
Backbone	ResNet-50	ViT-Small + CNN
Total clips	1,530	10,000
Test clips	306	1,500

The FAR reduction from 15.79% to 9.20% means the attack success rate drops from approximately 1-in-6 to 1-in-11. The EER improvement from 18.95% to 10.80% and AUC gain from 0.8736 to 0.9638 confirm that cross-modal audio-visual attention provides a genuine discriminative advantage over visual-only detection.

5.2 Data Scaling Experiment

Table 2 reports the data scaling experiment across three training corpus sizes, demonstrating a clear monotonic improvement trend confirming that transformer architectures benefit substantially from larger training corpora.

5.3 ROC Curves and Threshold Analysis

Figure 2 presents the ROC curve and FAR/FRR trade-off for AVConsist on the 1,500-sample LAV-DF test set. The AUC of 0.9638

Table 2: Data scaling experiment on LAV-DF. EER and AUC improve monotonically with training scale.

Total clips	Test clips	EER (%)	AUC
300	45	39.92	0.6957
1,000	150	16.67	0.8823
10,000	1,500	10.80	0.9638

indicates strong discrimination ability. The EER operating point occurs at threshold ≈ 0.25 , where FAR $\approx$ FRR $\approx 10.80\%$.




Figure 2: Left: ROC curve for AVConsist on LAV-DF (AUC = 0.9638, EER = 10.80%). Right: FAR/FRR trade-off. EER operating point marked at threshold ≈ 0.25 .

At the biometrically critical FAR $\leq 5\%$ operating point, AVConsist achieves FRR $\approx 18\%$, confirming a favorable security-utility trade-off for high-assurance deployment contexts.

5.4 Error Analysis

False negatives (fake accepted as real, FN = 97): Predominantly high-quality partial manipulations where the altered audio or video segment is short (< 1 second) and the unaltered majority of the clip dominates the frame-averaged features. The cross-modal attention module is unable to localize the manipulation without explicit temporal localization supervision, motivating future work on segment-level detection.

False positives (real rejected as fake, FP = 69): Occur in real clips with atypical audio characteristics—heavy background noise, reverb, or accented speech—that cause the CNN audio encoder to produce out-of-distribution embeddings flagged as synthesized. This motivates replacing the CNN voice encoder with the Audio Spectrogram Transformer (AST) [12] pretrained on AudioSet for more robust audio representations.

6 Code Repository

The GitHub repository contains the following additions for Checkpoint 2: <https://github.com/YashKhairnar/CS228-Project>

7 Conclusion and Future Work

AVConsist demonstrates that cross-modal audio-visual attention outperforms the visual-only ResNet-50 baseline on the challenging LAV-DF benchmark. The FAR reduction from 15.79% to 9.20% represents a concrete improvement in biometric security posture, and the AUC of 0.9638 confirms strong discriminative ability. The data

scaling experiment (EER: 39.92% $\rightarrow$ 16.67% $\rightarrow$ 10.80%) confirms that transformer architectures scale effectively with training data.

Planned improvements:

- (1) Replace CNN voice encoder with AST [12] pretrained on AudioSet for stronger audio representations.
- (2) Add temporal localization supervision for partial manipulation detection.
- (3) Evaluate cross-dataset generalization: train on FakeAVCeleb [10], test on LAV-DF zero-shot.
- (4) Apply adversarial training per the ART-AVDF [6] protocol.
- (5) Calibrated threshold selection using Bayesian cost-sensitive FAR/FRR calibration per ISO/IEC 19795-1 [8].

References

- [1] Y. Liu, H. Wang, and X. Chen. Cross-dataset generalization in deepfake detection: A systematic evaluation. 2023.
- [2] A. Rossler, D. Cozzolino, L. Verdoliva, C. Riess, J. Thies, and M. Nießner. Face-Forensics++: Learning to detect manipulated facial images. In *Proc. ICCV*, pages 1–11, 2019.
- [3] Y. Li, X. Yang, P. Sun, H. Qi, and S. Lyu. Frequency-aware discriminative feature learning supervised by single-center loss for face forgery detection. In *Proc. CVPR*, pages 1–10, 2021.
- [4] H. Feng, R. Hou, Y. Tang, and X. Hua. AVoid-DF: Audio-visual joint learning for detecting deepfake. *IEEE Trans. Information Forensics and Security*, 18:2969–2981, 2023.
- [5] Y. Zhou, L. Lim, C. Xu, and Y. Zhang. AVFF: Audio-visual feature fusion for video deepfake detection. In *Proc. CVPR*, pages 1–10, 2024.
- [6] Z. Cai, H. Guo, T. Chen, and W. Zhang. ART-AVDF: Adversarially robust training for audio-visual deepfake detection. In *Proc. ACM MM*, pages 1–10, 2024.
- [7] L. Zhang, B. Liu, Q. Chu, and N. Yu. Multimodal consistency-driven deepfake detection. In *Proc. ICIG*, 2025.
- [8] ISO/IEC. *ISO/IEC 19795-1:2021 — Biometric performance testing and reporting — Part 1: Principles and framework*. Standard, International Organization for Standardization, Geneva, Switzerland, 2021.
- [9] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *Proc. CVPR*, pages 770–778, 2016.
- [10] H. Khalid, S. Tariq, M. Kim, and S. S. Woo. FakeAVCeleb: A novel audio-video multimodal deepfake dataset. In *NeurIPS Datasets and Benchmarks Track*, 2021.
- [11] Z. Cai, K. Stefanov, A. Dhall, and M. Hayat. Do you really mean that? Content driven audio-visual deepfake dataset and multimodal method for temporal forgery localization. In *Proc. DICTA*, 2022.
- [12] Y. Gong, Y.-A. Chung, and J. Glass. AST: Audio spectrogram transformer. In *Proc. Interspeech*, pages 571–575, 2021.
- [13] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *Proc. ICLR*, 2021.